

Classification Task

1-Load the libraries:

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import confusion_matrix
```

2-Load the dataset:

```
In [3]: col_names = ["fLength", "fWidth", "fSize", "fConc", "fConc1", "fAsym", "fM3Long"]
magic_data = pd.read_csv("magic04.data", header=None, names=col_names) # read in
print(len(magic_data))
```

19020

3-Balance the Dataset:

```
In [16]: g_events = magic_data[magic_data['class'] == 'g']
h_events = magic_data[magic_data['class'] == 'h']
print("The length of gamma events =", len(g_events))
print("The length of the hadrons events =", len(h_events))

print("After the balancing:")
balanced_g = g_events.sample(n=len(h_events), random_state=42) #It balances the
print("The length of gamma events =", len(balanced_g))
balanced_data = pd.concat([balanced_g, h_events]).sample(frac=1, random_state=42)
print("The length of the total events =", len(balanced_data))
```

The length of gamma events = 12332

The length of the hadrons events = 6688

After the balancing:

The length of gamma events = 6688

The length of the total events = 13376

4-Encode the Labels:

```
In [5]: label_encoder = LabelEncoder() #It encodes the class labels as ML algos cannot w
balanced_data['class'] = label_encoder.fit_transform(balanced_data['class']) #It
```

5-Separate features and target variable:

```
In [6]: x = balanced_data.drop('class', axis=1) #It drops the class column and takes the
y = balanced_data['class'] # It takes the class column as the target variable
```

6-Split the data into training, validation and testing:

```
In [7]: x_train, x_temp = train_test_split(x, test_size=0.3, random_state=42, stratify=y)
y_train, y_temp = train_test_split(y, test_size=0.3, random_state=42, stratify=y)
```

```
x_val, x_test = train_test_split(x_temp, test_size=0.5, random_state=42, stratif
y_val, y_test = train_test_split(y_temp, test_size=0.5, random_state=42, stratif
print("The length of the training data =",len(x_train))
print("The length of the validation data =",len(x_val))
print("The length of the testing data =",len(x_test))
```

The length of the training data = 9363

The length of the validation data = 2006

The length of the testing data = 2007

7-Standardize the feature data (x_train,x_val,x_test):

```
In [8]: scaler = StandardScaler() #It scales the data by removing the mean and scaling t
x_train = scaler.fit_transform(x_train) #It fits the training data and then tran
x_val = scaler.transform(x_val) #It transforms the validation data
x_test = scaler.transform(x_test) #It transforms the testing data
print("The standardized training data",x_train)
print("The standardized validation data",x_val)
print("The standardized Testing data",x_test)
```

```

The standardized training data [[-0.67465295  0.03494   -0.12297294 ... -0.70774
943 -0.09676915
-0.97007624]
[-0.71885657 -0.30185083 -0.32184533 ...  0.72548882 -0.46537291
 0.76548842]
[-0.84905116 -0.60468813 -1.11838601 ...  0.27308091  0.0209516
 0.68958646]
...
[ 3.70106056  3.39132219  1.95509637 ... -1.81026703  2.17154426
 0.94663758]
[ 0.10815243  0.53609595  0.50706992 ...  0.55367925 -0.06791249
 1.00472735]
[ 2.38392593  0.02938362 -0.22325004 ...  0.74532998 -0.07173104
 0.76403464]]

The standardized validation data [[-0.65771323 -0.63338809 -0.50137707 ... -0.409
92857 -0.33840881
 0.64131963]
[-0.22172517 -0.32994037  0.42171877 ...  0.38791206 -1.00089986
 0.15473421]
[-0.66727821 -0.5987376  -0.31091366 ...  0.29210834 -1.03024732
 -0.12190625]
...
[ 1.54611287  1.19668548  1.73625265 ...  1.96495978 -0.96678572
 2.4096458 ]
[ 0.59984556  0.73806759  1.29814477 ... -1.42624273  0.05726689
 2.60902779]
[-0.16490473 -0.67235968 -0.2177842  ...  0.05446022  1.39422891
 1.05103564]]

The standardized Testing data [[-0.65349057 -0.61078765 -0.93086575 ... -0.443503
59  0.47984722
-0.83799573]
[-0.75948373 -0.55870858 -0.72883999 ...  0.22779337 -0.98878981
 0.36034146]
[ 0.3827771  -0.30599889  0.58989838 ...  0.48779309  0.14818712
 0.95425223]
...
[-0.85784838 -0.46313494 -0.61321651 ... -0.3881613  1.03430706
 -0.10417818]
[ 0.43075127 -0.26329742 -0.29346502 ... -0.63353985 -0.02827688
 -0.63062711]
[ 0.97528956  2.36657808  1.71649155 ... -3.11447714  0.66553226
 2.04945718]]

```

8-Create the KNN (K Neighbor Classifier):

```

In [9]: def evaluate_knn(k, x_train, y_train, x_val, y_val):
        knn_model = KNeighborsClassifier(n_neighbors=k) #It creates a KNN model with
        knn_model.fit(x_train, y_train) #It fits the model on the training data
        y_pred = knn_model.predict(x_val) #It predicts the class labels for the vali

        # Calculate performance metrics
        accuracy = accuracy_score(y_val, y_pred)
        precision = precision_score(y_val, y_pred)
        recall = recall_score(y_val, y_pred)
        f1 = f1_score(y_val, y_pred)
        cm = confusion_matrix(y_val, y_pred)

        return accuracy, precision, recall, f1, cm

```

9-Displaying the Matrices:

```
In [10]: def display_metrics(k, accuracy, precision, recall, f1, cm): #It displays the me
    print("\n" * 36)
    print(f"\n      **K = {k}**")
    print(f"Accuracy:  {accuracy:.4f}")
    print(f"Precision:  {precision:.4f}")
    print(f"Recall:      {recall:.4f}")
    print(f"F1 Score:   {f1:.4f}")
    print(f"Confusion Matrix:\n{cm}")
    print("\n" * 36)
```

10-Displaying the accuracy,precision,recall ,and f1 for each k value:

```
In [11]: def plot_metrics(k_values, metrics):

    # Extract the metrics for each k
    accuracies = [metrics[k]['accuracy'] for k in k_values]
    precisions = [metrics[k]['precision'] for k in k_values]
    recalls = [metrics[k]['recall'] for k in k_values]
    f1_scores = [metrics[k]['f1'] for k in k_values]

    # Plot the metrics
    plt.figure(figsize=(10, 6))
    plt.plot(k_values, accuracies, marker='o', label='Accuracy')
    plt.plot(k_values, precisions, marker='o', label='Precision')
    plt.plot(k_values, recalls, marker='o', label='Recall')
    plt.plot(k_values, f1_scores, marker='o', label='F1-Score')

    # Add labels and title
    plt.xlabel('k (Number of Neighbors)')
    plt.ylabel('Score')
    plt.title('K-NN Performance Metrics for Different k Values')
    plt.legend()
    plt.grid(True)
    plt.show()
```

11-Displaying the Confusion Matrix for all k values:

```
In [12]: def plot_confusion_matrix(cm, k):
    # Plot the confusion matrix
    plt.figure(figsize=(5, 5))
    plt.imshow(cm, interpolation='nearest', cmap=plt.cm.RdPu)
    plt.colorbar()
    plt.xticks([0, 1], ["g", "h"])
    plt.yticks([0, 1], ["g", "h"])
    plt.xlabel("Predicted Labels")
    plt.ylabel("True Labels")
    plt.title(f"Confusion Matrix (K = {k})")
    plt.show()
```

12-Evaluate the KNN Models:

```
In [13]: def test_knn(k_values, x_train, y_train, x_val, y_val):
    # Evaluate K-NN models for different k-values
    metrics = {}
    accuracy_scores = []
```

```

print("\n*** K-NN Performance Evaluation ***")
for k in k_values:
    # Evaluate K-NN model
    accuracy, precision, recall, f1, cm = evaluate_knn(k, x_train, y_train,

    # Store metrics
    metrics[k] = {
        'accuracy': accuracy,
        'precision': precision,
        'recall': recall,
        'f1': f1,
        'confusion_matrix': cm
    }

    # Display performance metrics
    display_metrics(k, accuracy, precision, recall, f1, cm)

    # Plot confusion matrix for the current k
    plot_confusion_matrix(cm, k)

    # Plot metrics for all k-values
    plot_metrics(k_values, metrics)

return metrics

```

```

In [14]: k_values = [1, 3, 5, 7, 9, 11, 13, 15] #It takes the values of k
         metrics = test_knn(k_values, x_train, y_train, x_val, y_val) #It tests the KNN m


```

```


*** K-NN Performance Evaluation ***
*****

**K = 1**
Accuracy: 0.7976
Precision: 0.8199
Recall: 0.7627
F1 Score: 0.7903
Confusion Matrix:
[[835 168]
 [238 765]]
*****

```



```
**K = 3**  
Accuracy: 0.8121  
Precision: 0.8533  
Recall: 0.7537  
F1 Score: 0.8004  
Confusion Matrix:  
[[873 130]  
 [247 756]]
```



****K = 5****

Accuracy: 0.8170

Precision: 0.8573


Recall: 0.7607

F1 Score: 0.8061

Confusion Matrix:

[[876 127]

[240 763]]



****K = 7****

Accuracy: 0.8180

Precision: 0.8709


Recall: 0.7468

F1 Score: 0.8041


Confusion Matrix:

[[892 111]

[254 749]]



```
**K = 9**  
Accuracy: 0.8161  
Precision: 0.8729  
Recall: 0.7398  
F1 Score: 0.8009  
Confusion Matrix:  
[[895 108]  
 [261 742]]
```



****K = 11****

Accuracy: 0.8220

Precision: 0.8791


Recall: 0.7468

F1 Score: 0.8075

Confusion Matrix:

[[900 103]

[254 749]]



****K = 13****

Accuracy: 0.8156

Precision: 0.8710


Recall: 0.7408

F1 Score: 0.8006

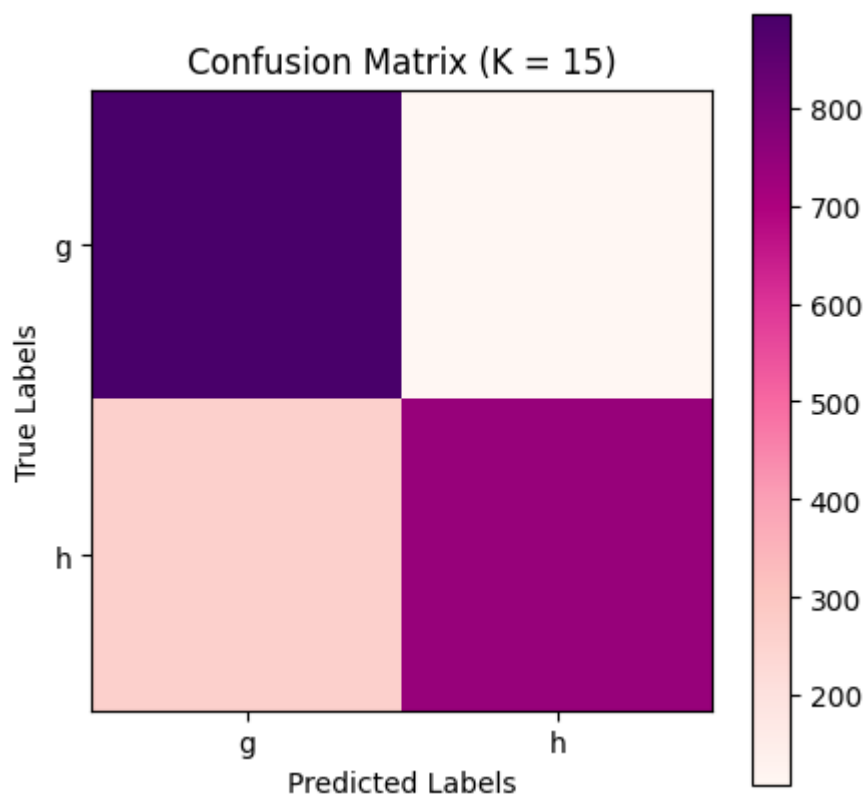
Confusion Matrix:


[[893 110]

[260 743]]



```
**K = 15**  
Accuracy: 0.8156  
Precision: 0.8728  
Recall: 0.7388  
F1 Score: 0.8002  
Confusion Matrix:  
[[895 108]  
 [262 741]]
```





13-The Best K value:

```
In [15]: best_k = max(metrics.keys(), key=lambda k: metrics[k]['accuracy']) #It selects the best k
best_accuracy = metrics[best_k]['accuracy'] #It selects the best accuracy based on best k
print(f"Best k value: {best_k} with accuracy: {best_accuracy:.4f}")
```

Best k value: 11 with accuracy: 0.8220

14-Comments:

for k=1:

- Accuracy (0.7976): This indicates that approximately 79.76% of the predictions made by the model are correct.
- Precision (0.8199): This indicates that when the model predicts a positive class (e.g., 'g'), it is correct about 81.99% of the time.
- Recall (0.7627): This indicates that the model correctly identifies 76.27% of all actual positive cases.
- F1 Score (0.7903): This is the harmonic mean of precision and recall, providing a single metric that balances both concerns. A score of 0.7903 suggests a good balance between precision and recall.

Confusion Matrix:

- True Positives (835): The number of correct positive predictions.
- False Positives (168): The number of incorrect positive predictions.
- False Negatives (238): The number of incorrect negative predictions.
- True Negatives (765): The number of correct negative predictions.

for k=3:

- Accuracy (0.8121): This indicates that approximately 81.21% of the predictions made by the model are correct.
- Precision (0.8533): This indicates that when the model predicts a positive class (e.g., 'g'), it is correct about 85.33% of the time.
- Recall (0.7537): This indicates that the model correctly identifies 75.37% of all actual positive cases.
- F1 Score (0.8004): This is the harmonic mean of precision and recall, providing a single metric that balances both concerns. A score of 0.8004 suggests a good balance between precision and recall.

Confusion Matrix:

- True Positives (873): The number of correct positive predictions.
- False Positives (130): The number of incorrect positive predictions.
- False Negatives (247): The number of incorrect negative predictions.
- True Negatives (756): The number of correct negative predictions.

for k=5:

- Accuracy (0.8170): This indicates that approximately 81.70% of the predictions made by the model are correct.
- Precision (0.8573): This indicates that when the model predicts a positive class (e.g., 'g'), it is correct about 85.73% of the time.
- Recall (0.7607): This indicates that the model correctly identifies 76.07% of all actual positive cases.
- F1 Score (0.8061): This is the harmonic mean of precision and recall, providing a single metric that balances both concerns. A score of 0.8061 suggests a good balance between precision and recall.

Confusion Matrix:

- True Positives (876): The number of correct positive predictions.
- False Positives (127): The number of incorrect positive predictions.
- False Negatives (240): The number of incorrect negative predictions.
- True Negatives (763): The number of correct negative predictions.

for k=7:

- Accuracy (0.8180): This indicates that approximately 81.80% of the predictions made by the model are correct.
- Precision (0.8709): This indicates that when the model predicts a positive class (e.g., 'g'), it is correct about 87.09% of the time.

- Recall (0.7468): This indicates that the model correctly identifies 74.68% of all actual positive cases.
- F1 Score (0.8041): This is the harmonic mean of precision and recall, providing a single metric that balances both concerns. A score of 0.8041 suggests a good balance between precision and recall.

Confusion Matrix:

- True Positives (892): The number of correct positive predictions.
- False Positives (111): The number of incorrect positive predictions.
- False Negatives (254): The number of incorrect negative predictions.
- True Negatives (749): The number of correct negative predictions.

for k=9:

- Accuracy (0.8161): This indicates that approximately 81.61% of the predictions made by the model are correct.
- Precision (0.8729): This indicates that when the model predicts a positive class (e.g., 'g'), it is correct about 87.29% of the time.
- Recall (0.7398): This indicates that the model correctly identifies 73.98% of all actual positive cases.
- F1 Score (0.8009): This is the harmonic mean of precision and recall, providing a single metric that balances both concerns. A score of 0.8009 suggests a good balance between precision and recall.

Confusion Matrix:

- True Positives (895): The number of correct positive predictions.
- False Positives (108): The number of incorrect positive predictions.
- False Negatives (261): The number of incorrect negative predictions.
- True Negatives (742): The number of correct negative predictions.

for k=11:

- Accuracy (0.8220): This indicates that approximately 82.20% of the predictions made by the model are correct.
- Precision (0.8791): This indicates that when the model predicts a positive class (e.g., 'g'), it is correct about 87.91% of the time.
- Recall (0.7468): This indicates that the model correctly identifies 74.68% of all actual positive cases.
- F1 Score (0.8075): This is the harmonic mean of precision and recall, providing a single metric that balances both concerns. A score of 0.8075 suggests a good balance between precision and recall.

Confusion Matrix:

- True Positives (900): The number of correct positive predictions.
- False Positives (103): The number of incorrect positive predictions.
- False Negatives (254): The number of incorrect negative predictions.
- True Negatives (749): The number of correct negative predictions.

for k=13:

- Accuracy (0.8156): This indicates that approximately 81.56% of the predictions made by the model are correct.
- Precision (0.8710): This indicates that when the model predicts a positive class (e.g., 'g'), it is correct about 87.10% of the time.
- Recall (0.7408): This indicates that the model correctly identifies 74.08% of all actual positive cases.
- F1 Score (0.8006): This is the harmonic mean of precision and recall, providing a single metric that balances both concerns. A score of 0.8006 suggests a good balance between precision and recall.

Confusion Matrix:

- True Positives (893): The number of correct positive predictions.
- False Positives (110): The number of incorrect positive predictions.
- False Negatives (260): The number of incorrect negative predictions.
- True Negatives (743): The number of correct negative predictions.

for k=15:

- Accuracy (0.8156): This indicates that approximately 81.56% of the predictions made by the model are correct.
- Precision (0.8728): This indicates that when the model predicts a positive class (e.g., 'g'), it is correct about 87.28% of the time.
- Recall (0.7388): This indicates that the model correctly identifies 73.88% of all actual positive cases.
- F1 Score (0.8002): This is the harmonic mean of precision and recall, providing a single metric that balances both concerns. A score of 0.8002 suggests a good balance between precision and recall.

Confusion Matrix:

- True Positives (895): The number of correct positive predictions.
- False Positives (108): The number of incorrect positive predictions.
- False Negatives (262): The number of incorrect negative predictions.
- True Negatives (741): The number of correct negative predictions.

15-Tradeoffs Between the models:

When choosing the optimal value of k for a K-Nearest Neighbors (K-NN) classifier, there are several tradeoffs to consider:

- **Bias-Variance Tradeoff:**
 - **Low k (e.g., $k = 1$):**
 - **Low Bias:** The model can capture more details and is more flexible.
 - **High Variance:** The model is more sensitive to noise in the training data, leading to overfitting.
 - **High k (e.g., $k = 15$):**
 - **High Bias:** The model is less flexible and may miss some details.
 - **Low Variance:** The model is more stable and less sensitive to noise, leading to underfitting.
- **Accuracy:** Generally, as k increases, the accuracy tends to stabilize. However, too high a value of k can lead to underfitting, where the model is too simple to capture the underlying patterns in the data.
- **Precision and Recall:**
 - **Precision:** Higher values of k tend to increase precision, as the model becomes more conservative in making positive predictions, reducing false positives.
 - **Recall:** Lower values of k tend to increase recall, as the model is more aggressive in making positive predictions, reducing false negatives.
- **F1 Score:** The F1 score balances precision and recall. The optimal k value should provide a good balance between these two metrics. In this case, $k = 11$ has the highest F1 score, indicating a good balance.
- **Confusion Matrix:** The confusion matrix provides insights into the types of errors the model is making. For example, with lower k values, you might see more false positives and false negatives, while higher k values might reduce these errors but increase the number of true negatives.

Specific Tradeoffs Observed:

- **$k = 1$:**
 - **Pros:** High recall (0.7627), capturing more true positives.
 - **Cons:** Lower precision (0.8199) and higher variance, leading to more false positives.
- **$k = 3$:**
 - **Pros:** Improved precision (0.8533) and balanced recall (0.7537).
 - **Cons:** Slightly lower recall compared to $k = 1$.
- **$k = 5$:**
 - **Pros:** Good balance with high precision (0.8573) and recall (0.7607).
 - **Cons:** Slightly lower accuracy compared to higher k values.

- **k = 11:**
 - **Pros:** Highest accuracy (0.8220), highest precision (0.8791), and highest F1 score (0.8075).
 - **Cons:** Slightly lower recall (0.7468) compared to k = 5.
- **k = 15:**
 - **Pros:** High precision (0.8728) and stable performance.
 - **Cons:** Lower recall (0.7388) and slightly lower F1 score (0.8002) compared to k = 11.

16- The Conclusion:

The choice of k involves balancing the tradeoffs between bias and variance, precision and recall, and overall model stability.

In this case, k = 11 provides the best overall performance, but the specific tradeoffs should be considered based on the application's requirements and the importance of different metrics.

In []: